

# 4-bit Quantization Model on Image Classification

**COMS 6998**

**Practical Deep Learning System Performance  
Final Project**

**Group:**

**Praditya Raudi Avinanto (pra2118)**

**Rifqi Luthfan (rl3154)**

**Date Submitted: Dec 16, 2021**



# Contents

1. **Executive summary**
2. **Problem motivation & Background work**
3. **Technical challenges**
4. **Approach & Solution diagram/architecture**
5. **Implementation details & Experiment design**
6. **Demo**
7. **Experimental evaluation**
8. **Conclusion**

# EXECUTIVE SUMMARY

# Executive Summary

1. Quantization is one of the techniques to reduce model size and computational complexity which can then be implemented in edge devices (Mobile Phones, IoT devices)
  - Currently, PyTorch and Tensorflow supports only 8-bit integer quantization
2. In this project, we explore converting a 32-bit float neural network (NN) model into a precision lower than 8-bit integer NN model
  - We experimented using 8,7,6,5,4 bits quantization for two models (ResNet-18 and ResNet-50) for two datasets (CIFAR10 and ImageNette)
  - We experimented both Post Training Quantization and Quantization Aware Training
3. We found the different effect of each model responding to different bitwidth quantization

# PROBLEM MOTIVATION & BACKGROUND WORK

# Problem Motivation

1. Neural networks have evolved without regard to model complexity and computational efficiency
  - Requires powerful computation machine with large storage capacity
  - **Not well suited for edge devices** with lower computation power
2. Quantization, a method to perform computations and storing tensors at lower bitwidths [2], can be utilized to solve this problem, however:
  - We usually see a significant drop in accuracy performance when moving to a lower precision bits
  - The current implementation of quantization only support 8 bits quantization





Image from [1] indicating higher accuracy implies higher computation & memory

## [1] [Benchmark Analysis of Representative Deep Neural Network Architectures](#)

## [2] [Quantization - PyTorch](#)

# What is Quantization?



1. Neural Networks models usually run arithmetic operations for the weights and activations in 32 bits floating point
2. Question arise: can neural networks run in lower bit precision?
  - This will allow model to be more compact
  - The concept behind quantization is mapping the floating point weights and activations into a different integer bit scheme
    - i. Get zero\_point and scale of the float32 values
    - ii. Do quantize/dequantize based on the formula
$$\text{float\_value} = (\text{int\_value} - \text{zero\_point}) * \text{scale}$$
$$\text{int\_value} = \text{float\_value} / \text{scale} - \text{zero\_point}$$

[3] [Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#)

# Types of Quantization

## 1. Post Training Quantization (PTQ)

- Model training is done using normal float32 weights and activations, and then quantization is applied after the training is completed
  - i. Simple to apply
  - ii. Accuracy downgrade might be big

## 2. Quantization Aware Training (QAT)

- Simulates the quantization of weights and activations during training's forward pass process
  - i. More robust accuracy performance because quantization error is taken into account
  - ii. More involved than other option (involving fine-tuning and access to full dataset)

[3] [Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#)



# TECHNICAL CHALLENGES

# Frameworks only support 8 bit quantization

## Quantization Operation coverage

Quantized Tensors support a limited subset of data manipulation methods of the regular full-precision tensor. For NN operators included in PyTorch, we restrict support to:

1. 8 bit weights (`data_type = qint8`)
2. 8 bit activations (`data_type = quint8`)

Source: <https://pytorch.org/docs/stable/quantization.html>



Source: [https://www.tensorflow.org/lite/performance/post\\_training\\_quantization](https://www.tensorflow.org/lite/performance/post_training_quantization)

## Optimization Methods

There are several post-training quantization options to choose from. Here is a summary table of the choices and the benefits they provide:

| Technique                  | Benefits                     | Hardware                        |
|----------------------------|------------------------------|---------------------------------|
| Dynamic range quantization | 4x smaller, 2x-3x speedup    | CPU                             |
| Full integer quantization  | 4x smaller, 3x+ speedup      | CPU, Edge TPU, Microcontrollers |
| Float16 quantization       | 2x smaller, GPU acceleration | CPU, GPU                        |



# Quantization affects all Neural Networks operations


All of the operations are executed on integer tensors rather than floating point values on Quantized model

- We need to implement custom modules and functions that supports quantization lower than 8 bits from vanilla framework


[3] [Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#)

# APPROACH & SOLUTION DIAGRAM / ARCHITECTURE

# Neural Network vs Quantized Neural Network



Source: <https://medium.com/analytics-vidhya/neural-network-part1-inside-a-single-neuron-fee5e44f1e>



$Q_x = \{x\_scale, x\_zero\_point, x\_max, x\_min, q\_max, q\_min\}$   
 $Q_w = \{w\_scale, w\_zero\_point, w\_max, w\_min, q\_max, q\_min\}$   
 $Q_a = \{a\_scale, a\_zero\_point, a\_max, a\_min, q\_max, q\_min\}$

## QUANT & DEQUANT OP


$$\text{Scale} = (v\_max - v\_min) / (q\_max - q\_min)$$

$$\text{Zero\_point} = q\_max - v\_max / \text{scale}$$

$$\text{float\_value} = (\text{int\_value} - \text{zero\_point}) * \text{scale}$$

$$\text{int\_value} = \text{float\_value} / \text{scale} - \text{zero\_point}$$

# Post Training Quantization



$Q_x = \{x\_scale, x\_zero\_point, x\_max, x\_min, q\_max, q\_min\}$   
 $Q_w = \{w\_scale, w\_zero\_point, w\_max, w\_min, q\_max, q\_min\}$   
 $Q_a = \{a\_scale, a\_zero\_point, a\_max, a\_min, q\_max, q\_min\}$

## QUANT & DEQUANT OP

Scale =  $(v\_max - v\_min) / (q\_max - q\_min)$

Zero\_point =  $q\_max - v\_max / scale$

float\_value =  $(int\_value - zero\_point) * scale$

int\_value =  $float\_value / scale - zero\_point$




Image from [3] [Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#)

# Quantization Aware Training

1. We implemented “Fake” Quantization Function that quantize and dequantize the weights and activations of the network during forward pass to simulate precision loss.
2. We also implemented QParam as a module responsible for updating the quantization parameters during training and also quantize the tensor if needed.

```
class FakeQuantize(Function):  
    @staticmethod  
    def forward(ctx, x, qparam):  
        x = qparam.quantize_tensor(x)  
        x = qparam.dequantize_tensor(x)  
        return x  
  
    @staticmethod  
    def backward(ctx, grad_output):  
        return grad_output, None
```

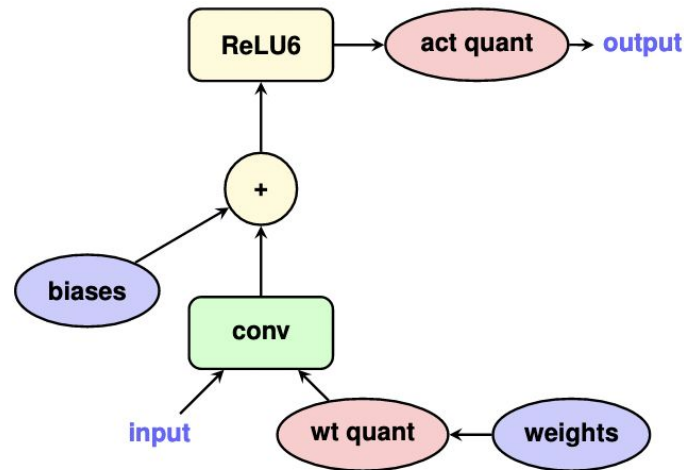
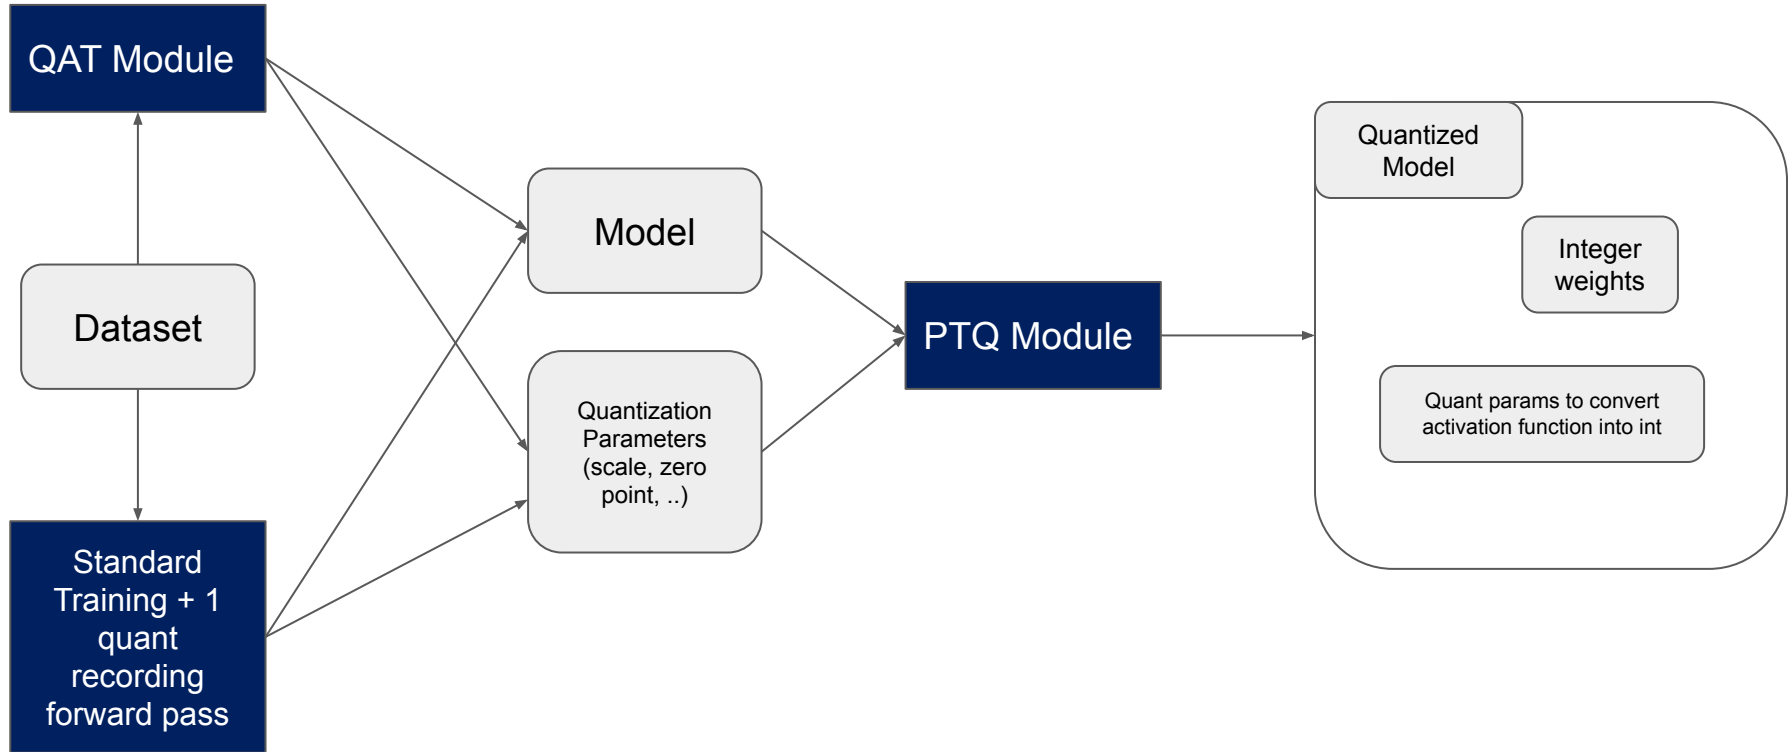


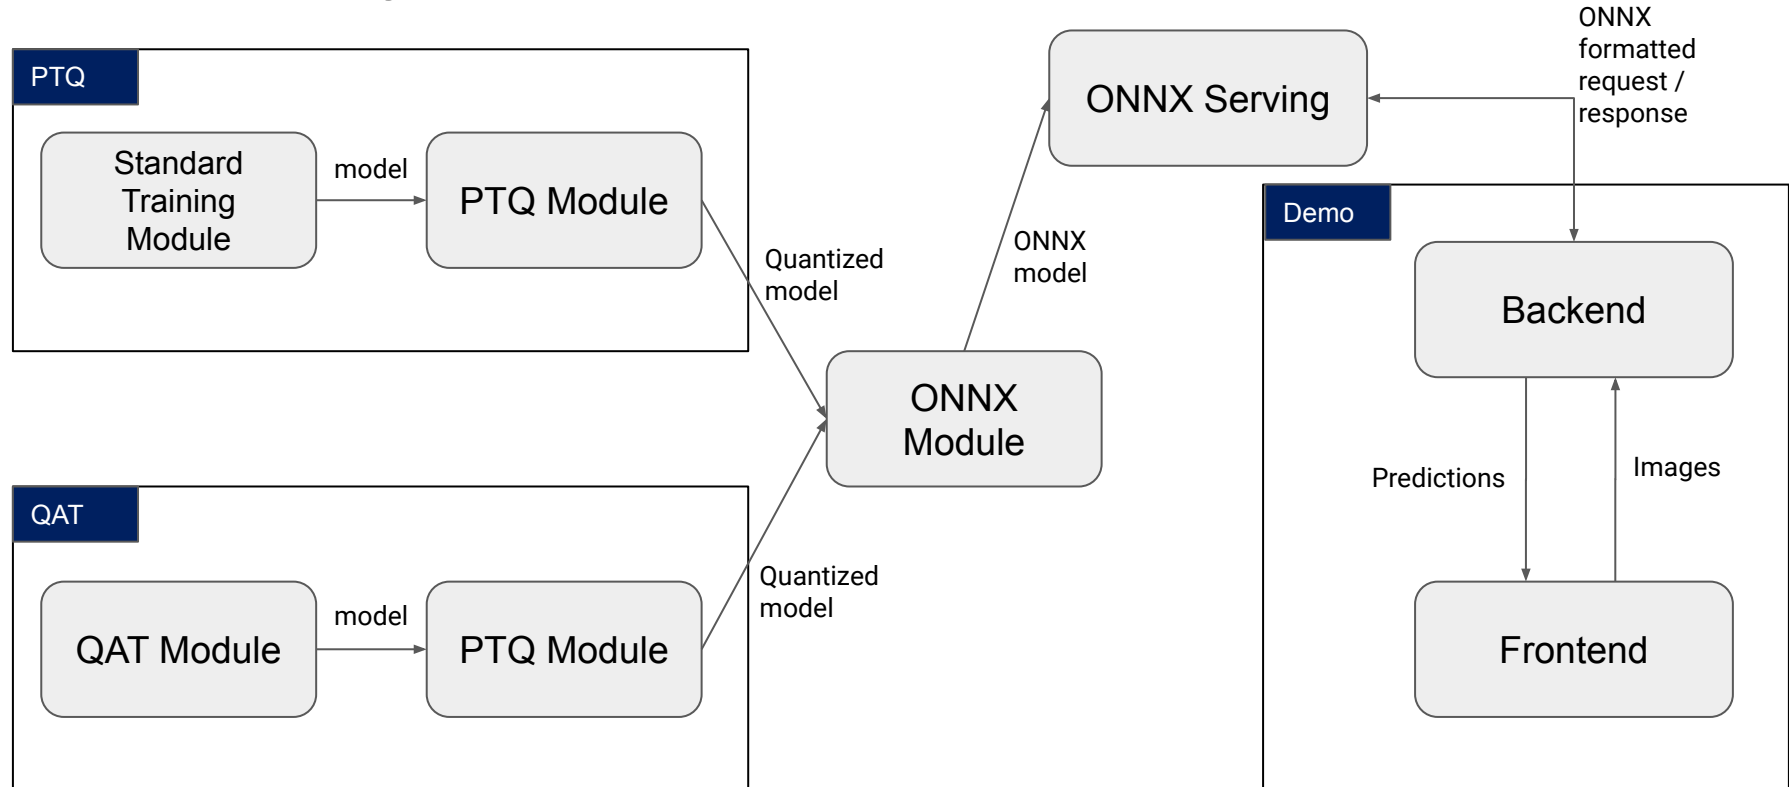
Image from [3] [Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#)

# Quantization Flow





# Solution Diagram / Architecture



# IMPLEMENTATION DETAILS & EXPERIMENT DESIGN

# Experiment Design

## Datasets:

- Imagenette
  - subset of 10 easily classified classes from Imagenet

- CIFAR10

## Model Architectures (pretrained on Imagenet):

- Resnet18
- Resnet50

## Quantization Methods:

- Post Training Quantization
- Quantization Aware Training

## Hardware:

- NVIDIA Tesla V100
- 8 vCPU 30gb RAM (n1-standard-8 GCP)

## Bitwidths:

- Full precision: 32 bit
- Low precision: 8 bit, 7 bit, 6 bit, 5 bit, 4 bit

## Metrics:

- Accuracy
- Model Size
- Inference Time

## Framework:

- Pytorch
- Fastai (for dataloader)
- ONNX (for inference runtime)

# Implementation Details

1. Implement quantization module from scratch
  - Use Pytorch module register buffer to keep quantization parameters
    - i. zero point, scale, vmin, vmax, and num bit
  - Implement quantization-compatible layer for conv, batchnorm, relu, maxpool, average pool, linear, and add layer
  - Conv, BN, and relu fusion function implementation to speedup runtime
  - These layers is used to implement quantization-compatible basic and bottleneck block for Resnet
2. QAT implementation
  - Fake quantization Function
  - Quantization parameters updates on training
3. PTQ implementation
  - Quantization / model conversion.

# Custom Layers and Models

1. We implemented custom Convolutional, Batch Normalization, Relu, Linear, and Addition Layers differently for PTQ and QAT. All PTQ weights & activation outputs are already in integers while QAT has FakeQuant and QParam.
2. Those layers are then used to create a custom ResNet Model then we can do Post Training Quantization and Quantization Aware Training

```
class QConvBNReLU(nn.Module):
    def __init__(self, in_channel, out_channel, kernel_size, relu, stride=1, padding=0, dilation=1):
        super(QConvBNReLU, self).__init__()
        self.conv = nn.Conv2d(in_channel, out_channel, kernel_size, stride, padding, dilation)
        self.bn = nn.BatchNorm2d(out_channel)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.conv(x)
        x = self.bn(x)
        x = self.relu(x)
        return x
```

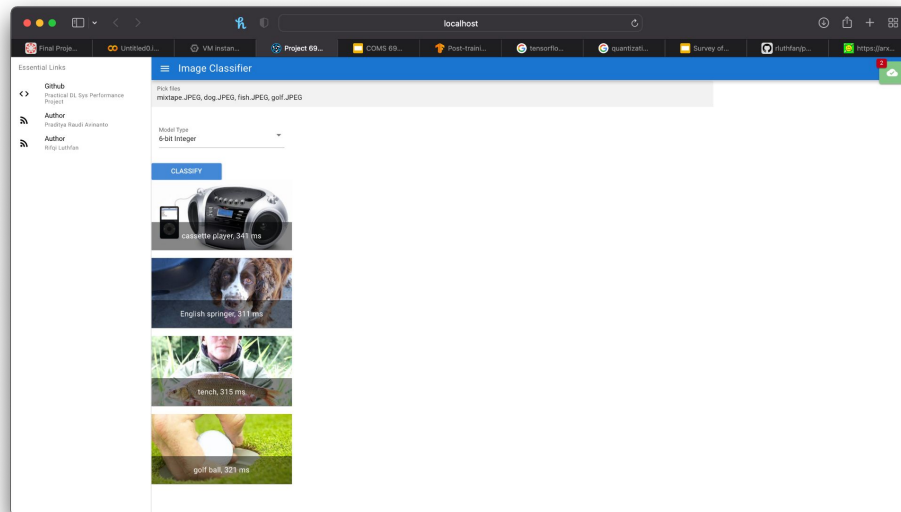
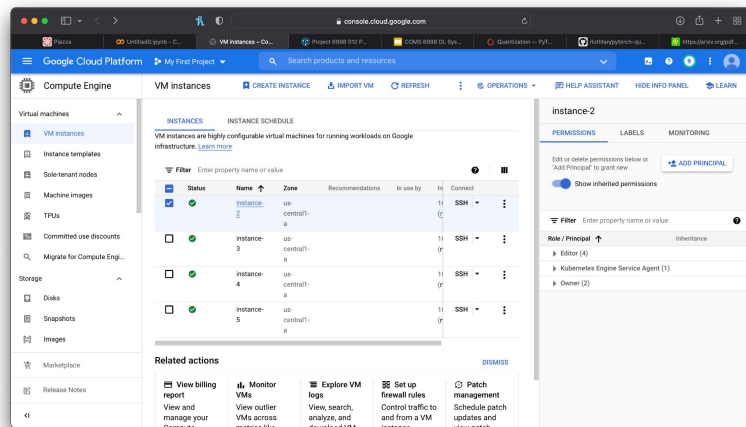
```
class CConvBNReLU2d(nn.Module):
    def __init__(self, in_channel, out_channel, kernel_size, stride=1, padding=0, bias=False, dilation=1,
                 q_num_bit=8, start=False, affine=True, relu=False, state_dict_names=[], qat=False):
        super(CConvBNReLU2d, self).__init__()
        self.conv = nn.Conv2d(in_channel, out_channel, kernel_size, stride, padding, dilation)
        self.bn = nn.BatchNorm2d(out_channel)
        self.relu = nn.ReLU()
        self.qat = qat
        self.q_num_bit = q_num_bit
        self.start = start
        self.affine = affine
        self.state_dict_names = state_dict_names

    def forward(self, x):
        if not self.qat:
            x = F.conv2d(x, self.conv.weight, self.conv.bias, self.conv.stride, self.conv.padding, self.conv.dilation)
            x = F.batch_norm(x, self.bn.running_mean, self.bn.running_var, self.bn.weight, self.bn.bias)
            x = self.relu(x)
            return x
        else:
            if self.start:
                self.q_in.quantize_update(x)
                if self.qat:
                    x = FakeQuantize.apply(x, self.q_in)
            weight, bias = fold_bn(self.conv.weight, self.conv.bias, self.bn.running_mean, self.bn.running_var, self.bn.weight, self.bn.bias)
            if self.q_w:
                self.q_w.quantize_update(weight)
                if self.qat:
                    weight = FakeQuantize.apply(weight, self.q_w)
            if self.bias:
                self.q_b.scale = self.q_w.scale
                self.q_b.zero_point = torch.tensor(0)
                if self.qat:
                    bias = FakeQuantize.apply(bias, self.q_b)
```

Sample  
custom layer

# Web Application

1. Four ONNXRuntime servers  
(Resnet50 32, 8, 6, 4 bit), 1 vCPU,  
3.75 GB RAM each
2. Backend Flask
3. Frontend VueJS



# DEMO

# EXPERIMENT EVALUATION



# Validation Accuracy of Quantized model


| quantization | dataset  |  | cifar10  |          | imagenette |          |
|--------------|----------|--|----------|----------|------------|----------|
|              | model    |  | resnet18 | resnet50 | resnet18   | resnet50 |
|              | num_bits |  |          |          |            |          |
| base         | 32       |  | 0.8039   | 0.8460   | 0.9465     | 0.9603   |
| ptq          | 8        |  | 0.7968   | 0.2418   | 0.9439     | 0.8089   |
|              | 7        |  | 0.7926   | 0.2377   | 0.9404     | 0.7911   |
|              | 6        |  | 0.7539   | 0.1806   | 0.9027     | 0.7554   |
|              | 5        |  | 0.5153   | 0.1378   | 0.2550     | 0.3654   |
|              | 4        |  | 0.1336   | 0.1061   | 0.1284     | 0.0721   |

| quantization | dataset  |  | cifar10  |          | imagenette |          |
|--------------|----------|--|----------|----------|------------|----------|
|              | model    |  | resnet18 | resnet50 | resnet18   | resnet50 |
|              | num_bits |  |          |          |            |          |
| base         | 32       |  | 0.8039   | 0.8460   | 0.9465     | 0.9603   |
| qat          | 8        |  | 0.8057   | 0.7934   | 0.9439     | 0.9513   |
|              | 7        |  | 0.8118   | 0.7908   | 0.9401     | 0.9442   |
|              | 6        |  | 0.7689   | 0.7334   | 0.7689     | 0.8433   |
|              | 5        |  | 0.6344   | 0.4095   | 0.8061     | 0.4438   |
|              | 4        |  | 0.2257   | 0.0990   | 0.1332     | 0.1129   |


1. Post Training Quantization (left) will benefit more from a similar source dataset
2. Quantization Aware Training (right) has lower accuracy drop compared to Post Training Quantization (left)
  - Quantization error is taken into account during learning process
3. Even with QAT, the more complex the model is (ResNet-50), the worse the resulting accuracy drop when doing lower bitwidths quantization (true for both dataset)
  - This is possibly caused by the more complex float32 values required in deeper model

# Model lose capacity to learn with 4 bits quantization

IMAGENETTE\_RESNET18 Metrics



IMAGENETTE\_RESNET50 Metrics



# Model Size

|        | Resnet18 | Resnet50 |
|--------|----------|----------|
| 32 bit | 44.9 MB  | 94.8 MB  |
| 8 bit  | 11.2 MB  | 23.7 MB  |
| 7 bit  | 11.2 MB  | 23.7 MB  |
| 6 bit  | 11.2 MB  | 23.7 MB  |
| 5 bit  | 11.2 MB  | 23.7 MB  |
| 4 bit  | 11.2 MB  | 23.7 MB  |

Most current frameworks still only support 8 bit integers

# Web Inference Time

Image Classifier

Pick files  
mixtape.JPEG, dog.JPEG, fish.JPEG, golf.JPEG

Model Type  
32-bit Float

CLASSIFY

cassette player, 289 ms

English springer, 285 ms

tench, 280 ms

golf ball, 289 ms

Image Classifier

Pick files  
mixtape.JPEG, dog.JPEG, fish.JPEG, golf.JPEG

Model Type  
8-bit Integer

CLASSIFY

cassette player, 325 ms

English springer, 334 ms

tench, 314 ms

golf ball, 313 ms

Image Classifier

Pick files  
mixtape.JPEG, dog.JPEG, fish.JPEG, golf.JPEG

Model Type  
6-bit Integer

CLASSIFY

cassette player, 341 ms

English springer, 311 ms

tench, 315 ms

golf ball, 321 ms

Image Classifier

Pick files  
mixtape.JPEG, dog.JPEG, fish.JPEG, golf.JPEG

Model Type  
4-bit Integer

CLASSIFY

cassette player, 326 ms

English springer, 368 ms

tench, 322 ms

tench, 301 ms

# CONCLUSION

# Conclusion

1. Quantization Aware Training performs better than Post Training Quantization since it incorporates quantization loss during the training.
2. Compressing the model to lower precision than 8-bit is possible and the model can still achieve acceptable accuracy on 5-bit precision.
3. Model is unable to learn or losing the capacity when trained on 4-bit precision.
4. For x86-based processor, 4-bit, 5-bit, 6-bit, 7-bit and 8-bit quantization don't make any difference in size and memory usage. Empirically, the lowest possible representation of integer in python is 8-bit.
5. ONNX can't leverage PyTorch quantization.
  - Theoretically, the advantage of 4-bit quantization is 2x reduction in bandwidth, which can increase throughput even though model size is the same as 8 bit